

Running Word2Vec with Chinese/English Wikipedia Dump

LUO Linkai

Deep Learning Research & Application Center

March 6, 2017

Outline

Word2Vec Revisited

English Wikipedia Word2Vec

Chinese Wikipedia Word2Vec

Concluding Remarks

Word2Vec Revisited

What is Word2Vec

- The corpus: “The cat sat on the mat”
- One-hot representation: “sat” = $[0, 0, 0, 1]^T$
- Distributed representation: “sat” = $[0.01, -0.02, 0.03]^T$
- Word similarity: “sat” \approx “sit”
- Word2Vec is the technique to obtain such meaningful word vectors given corpus

Idea


<start> The cat sat on the mat <end>

<start> The cat sat on the mat <end>


<start> The cat sat on the mat <end>

- CBOW: $p(\text{sat} | [\text{the}, \text{cat}, \text{on}, \text{the}])$ (treats an entire context as one observation)
- Skip-gram: $p(\text{the} | \text{sat}) p(\text{cat} | \text{sat}) p(\text{on} | \text{sat}) p(\text{the} | \text{sat})$ (treats each context-target pair as a new observation)

Word2Vec Algorithms

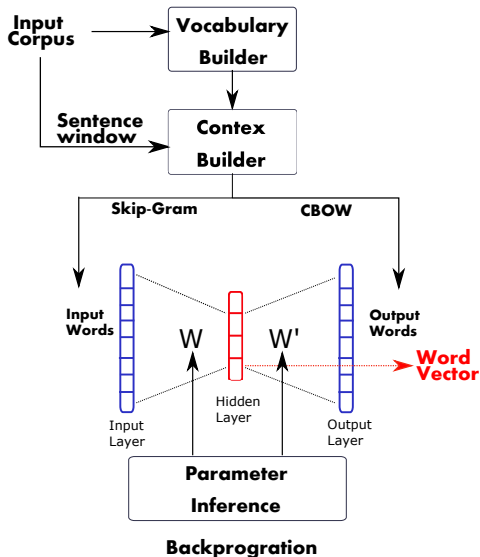


CBOW

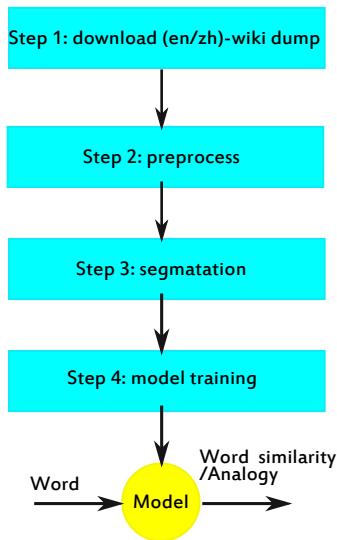


Skip-gram

Word2Vec Decomposed



Steps of Wikipedia Word2Vec



English Wikipedia Word2Vec

Download English Wikipedia Dump

- `https://dumps.wikimedia.org/enwiki/latest/enwiki-latest-pages-articles.xml.bz2`
- 13+ G file size, 4,000,000+ articles
- Preprocess, 2G+ memory required

```
1 # Remove words occurring less than 20 times, and words
   occurring in more than 10% of the documents. (keep_n is
   the vocabulary size)
2 wiki.dictionary.filter_extremes(no_below=20, no_above
   =0.1, ^^lkeep_n=100000)
```

Train Word2Vec Model

- Use gensim with the following settings:

```
1 sentences = SentencesIterator(wiki)
2 model = gensim.models.Word2Vec(sentences=sentences, size
    =200, min_count=3, window=5, workers=cores)
```

- Required RAM: $\mathcal{O}(\text{size} \cdot |V|)$
- ~ 24 hr training on virtual server

Word2Vec Result

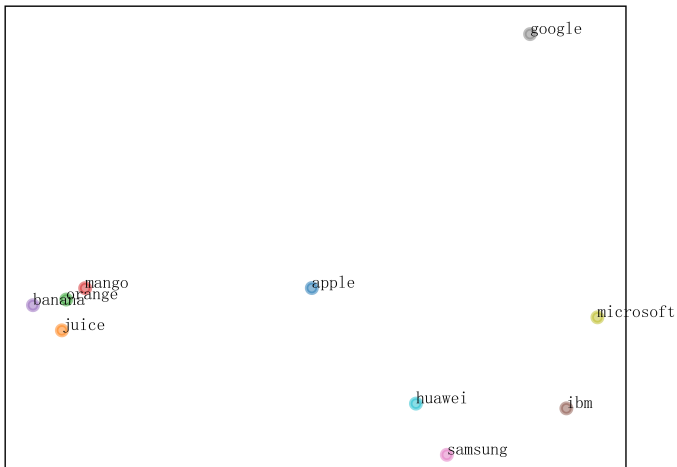
- Word similarity:

```
1 model = gensim.models.Word2Vec.load(os.path.join(MODEL_PATH
, 'word2vec.model'), mmap='r') # load large array
2 print(model.similarity('apple', 'mango'), model.similarity(
'apple', 'ibm'))
3 0.465574139702 0.531822866913
```

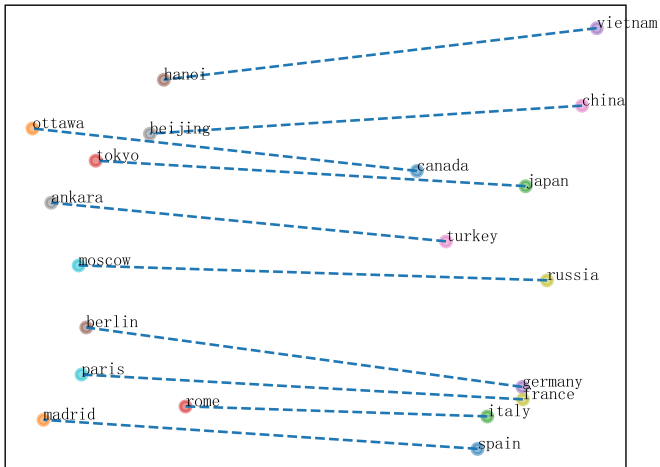
- Word Analogy:

```
1 King-man+woman
2
3 "queen" — similarity: 0.719937
4 "monarch" — similarity: 0.610968
5 "princess" — similarity: 0.608749
6 "prince" — similarity: 0.607932
7 "empress" — similarity: 0.586417
8 "throne" — similarity: 0.568553
9 "emperor" — similarity: 0.553523
10 "isabeau" — similarity: 0.540711
11 "amalasuntha" — similarity: 0.537015
```

Word Vector: “Apple”



Word Analogy: Country-Captain Paris



Chinese Wikipedia Word2Vec

Download Chinese Wikipedia Dump

- `https://dumps.wikimedia.org/zhwiki/latest/zhwiki-latest-pages-articles.xml.bz2`
- it contains traditional Chinese and simplified Chinese articles
- 1.3+ G file size, 230,000+ articles
- Preprocess, 2G+ memory required

Preprocessing

- Use **OpenCC** to translate from simplified Chinese to traditional Chinese
- Developed on C++, supporting Python

```
1 openCC = OpenCC('s2twp') # convert from Simplified Chinese  
   to Traditional Chinese  
2 converted = openCC.convert(original)  
3 print(openCC.convert('这是简体中文'))  
4 這是簡體中文  
5 print(openCC.convert('智能手机'))  
6 智慧手機  
7 print(openCC.convert('一条短信'))  
8 一條簡訊
```


Segmentation

- English uses whitespace, dots, etc to separate words, but Chinese does not
- 下雨天留客天留我不留
 - 下雨天留客, 天留我不留
 - 下雨天, 留客天, 留我不? 留
 - ...
- Many existing tools: stanford, fudan, jieba ...

Segmentation

```
1 import jieba
2 print('/'.join(jieba.cut('下雨天留客天留我不留',cut_all=False)))
3 下雨天/留客/天留/我/不留
4 '我畢業於香港科技大學化工系'
5 我/畢業/於/香港科技大學/化工系
6 '見了他,她變得很低很低,低到塵埃里,但她心里是歡喜的,從塵埃里開出
   花來'
7 見/了/他/,/她/變得/很/低/很/低/,/低到/塵埃/里/,/但/她/心里/是/歡
   喜/的/,/從/塵埃/里/開出/花來
```

File Format



Train Chinese Word2Vec Model

- Use gensim with the following settings:


```
1 sentences = SentencesIterator(wiki)
2 model = gensim.models.Word2Vec(sentences=sentences, size
    =300, min_count=3, window=5, workers=cores)
```

- ~ 12 hr training


Chinese Word2Vec Result

```
1 model = gensim.models.Word2Vec.load(os.path.join(MODEL_PATH, '
    word2vec.model'), mmap='r')
2 print(model.similarity('蘇軾', '東坡'))
3 0.612037855344
4
5 pprint(model.most_similar('奧斯卡'))
6 [('奧斯卡金像獎', 0.6850196123123169),
7 ('奧斯卡獎', 0.6607037782669067),
8 ('金球獎', 0.6433079242706299),
9 ('東尼獎', 0.5794373750686646),
10 ('金棕櫚獎', 0.5717018842697144),
11 ('艾美獎', 0.5554369688034058),
12 ('斷背山', 0.5518748760223389),
13 ('泰坦尼克號', 0.5482996702194214),
14 ('布克獎', 0.546217679977417),
15 ('學院獎', 0.5445188879966736)]
```

Word Similarity: Names



Word Analogy: Province-Capital Paris



Concluding Remarks

- Running Word2Vec with English/Chinese Wikipedia dump have acceptable performance in terms of word similarity/analogy
- Judge of Word2Vec's performance should also take machine translation into consideration
- Next ...
 - sentence pairs
 - constructing basic LSTMs